

Predictive Modeling for Business Forecasting

Feature Selection and Engineering:

```
import pandas as pd
```

```
# Load the dataset

file_path = '/mnt/data/Financial Distress.csv'

data = pd.read_csv(file_path)

# Display the first few rows of the dataset to understand its structure and
contents


data.head()
```

Company	Time	Financial Distress	x1	x2	x3	x4	\
0	1	1	0.010636	1.2810	0.022934	0.87454	1.21640
1	1	2	-0.455970	1.2700	0.006454	0.82067	1.00490
2	1	3	-0.325390	1.0529	-0.059379	0.92242	0.72926
3	1	4	-0.566570	1.1131	-0.015229	0.85888	0.80974
4	2	1	1.357300	1.0623	0.107020	0.81460	0.83593

	x5	x6	x7	...	x74	x75	x76	x77	x78	\
0	0.060940	0.188270	0.52510	...	85.437	27.07	26.102	16.000	16.0	
1	-0.014080	0.181040	0.62288	...	107.090	31.31	30.194	17.000	16.0	
2	0.020476	0.044865	0.43292	...	120.870	36.07	35.273	17.000	15.0	
3	0.076037	0.091033	0.67546	...	54.806	39.80	38.377	17.167	16.0	
4	0.199960	0.047800	0.74200	...	85.437	27.07	26.102	16.000	16.0	

	x79	x80	x81	x82	x83
0	0.2	22	0.060390	30	49
1	0.4	22	0.010636	31	50
2	-0.2	22	-0.455970	32	51
3	5.6	22	-0.325390	33	52
4	0.2	29	1.251000	7	27

[5 rows x 86 columns]



Model Selection and Development:

```
import seaborn as sns
import matplotlib.pyplot as plt
```

```
# Calculate correlations between 'Financial Distress' and other features
correlations = data.corr()['Financial Distress'].sort_values(ascending=False)
```

```
# Plot the correlations
plt.figure(figsize=(10, 15))
sns.barplot(y=correlations.index, x=correlations.values, palette="vlag")
```

```
plt.title('Correlation of Features with Financial Distress')
plt.show()
```

```
# Display correlations for reference
```

Correlations

```
Financial Distress    1.000000
x81                   0.442916
x48                   0.435318
x25                   0.376786
x10                   0.338298
```

...

```
x62                   -0.143793
x63                   -0.150362
x50                   -0.155569
x61                   -0.157380
x64                   -0.173126
```

```
Name: Financial Distress, Length: 86, dtype: float64
```

```
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
```

```
# Selecting a subset of features based on correlation (both positive and
negative)
selected_features = correlations[abs(correlations) > 0.15].index.tolist()
selected_features.remove('Financial Distress')    # Removing the target
variable
```

```
# Preparing the data with the selected features
```

```
X = data[selected_features]
y = data['Financial Distress']
```

```
# Splitting the data into training and testing sets
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42)
```

```
# Scaling the features
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Displaying the selected features for reference
selected_features
```

```
['x81',
 'x48',
 'x25',
 'x10',
 'x5',
 'x2',
 'x53',
 'x9',
 'x51',
 'x36',
 'x54',
 'x26',
 'x23',
 'x63',
 'x50',
 'x61',
 'x64']
```

Model Evaluation and Validation:

```
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, mean_absolute_error
import numpy as np
```

```
# Initialize and train the Linear Regression model
lr_model = LinearRegression()
lr_model.fit(X_train_scaled, y_train)
```

```
# Predict on the training set
y_train_pred = lr_model.predict(X_train_scaled)
```

```
# Evaluate the model on the training set

rmse_train = np.sqrt(mean_squared_error(y_train, y_train_pred))
mae_train = mean_absolute_error(y_train, y_train_pred)

rmse_train, mae_train
```

(2.2691174452745315, 0.9171223258760326)

```
# Predict on the test set
y_test_pred = lr_model.predict(X_test_scaled)
```

```
# Evaluate the model on the test set

rmse_test = np.sqrt(mean_squared_error(y_test, y_test_pred))
mae_test = mean_absolute_error(y_test, y_test_pred)

rmse_test, mae_test

from sklearn.model_selection import cross_val_score
```

(1.6350901089837295, 0.9461840047357485)

```
# Perform cross-validation on the Linear Regression model

cv_rmse_scores = -cross_val_score(lr_model, X_train_scaled, y_train,
scoring='neg_root_mean_squared_error', cv=5)

cv_mae_scores = -cross_val_score(lr_model, X_train_scaled, y_train,
scoring='neg_mean_absolute_error', cv=5)
```

```
# Calculate the mean and standard deviation of the cross-validation scores

cv_rmse_mean, cv_rmse_std = cv_rmse_scores.mean(), cv_rmse_scores.std()
cv_mae_mean, cv_mae_std = cv_mae_scores.mean(), cv_mae_scores.std()

cv_rmse_mean, cv_rmse_std, cv_mae_mean, cv_mae_std
```

(3.7307861443638877,
2.6555682933004636,
1.0122196667350736,

0.07320137753658636)

```
from sklearn.ensemble import RandomForestRegressor
```

```
# Initialize and train the Random Forest Regressor model
rf_model = RandomForestRegressor(random_state=42, n_estimators=100)
rf_model.fit(X_train_scaled, y_train)
```

```
# Predict and evaluate on the training set
y_train_pred_rf = rf_model.predict(X_train_scaled)
rmse_train_rf = np.sqrt(mean_squared_error(y_train, y_train_pred_rf))
mae_train_rf = mean_absolute_error(y_train, y_train_pred_rf)
```

```
# Predict and evaluate on the test set
y_test_pred_rf = rf_model.predict(X_test_scaled)
rmse_test_rf = np.sqrt(mean_squared_error(y_test, y_test_pred_rf))
mae_test_rf = mean_absolute_error(y_test, y_test_pred_rf)

rmse_train_rf, mae_train_rf, rmse_test_rf, mae_test_rf
```

```
(1.0652199404546085,
 0.24280903619863814,
 1.1622058774860664,
 0.6274488595595917)
```

```
from sklearn.ensemble import GradientBoostingRegressor
```

```
# Initialize and train the Gradient Boosting Regressor model
gbr_model = GradientBoostingRegressor(random_state=42)
gbr_model.fit(X_train_scaled, y_train)
```

```
# Predict and evaluate on the training set
y_train_pred_gbr = gbr_model.predict(X_train_scaled)
rmse_train_gbr = np.sqrt(mean_squared_error(y_train, y_train_pred_gbr))
```

```
mae_train_gbr = mean_absolute_error(y_train, y_train_pred_gbr)
```

```
# Predict and evaluate on the test set

y_test_pred_gbr = gbr_model.predict(X_test_scaled)

rmse_test_gbr = np.sqrt(mean_squared_error(y_test, y_test_pred_gbr))

mae_test_gbr = mean_absolute_error(y_test, y_test_pred_gbr)


rmse_train_gbr, mae_train_gbr, rmse_test_gbr, mae_test_gbr

from sklearn.svm import SVR
```

```
(0.7804365959553509,          0.4660935530156846,          1.101318683759754,
0.6130144822746689)
```

```
# Initialize and train the Support Vector Regressor model

svr_model = SVR(kernel='rbf')

svr_model.fit(X_train_scaled, y_train)
```

```
# Predict and evaluate on the training set

y_train_pred_svr = svr_model.predict(X_train_scaled)

rmse_train_svr = np.sqrt(mean_squared_error(y_train, y_train_pred_svr))

mae_train_svr = mean_absolute_error(y_train, y_train_pred_svr)
```

```
# Predict and evaluate on the test set

y_test_pred_svr = svr_model.predict(X_test_scaled)

rmse_test_svr = np.sqrt(mean_squared_error(y_test, y_test_pred_svr))

mae_test_svr = mean_absolute_error(y_test, y_test_pred_svr)


rmse_train_svr, mae_train_svr, rmse_test_svr, mae_test_svr
```

```
(2.554341762607816,          0.5526126578022971,          1.067986146331273,
0.5814032646271093)
```

```
# Train the Decision Tree Regressor with default parameters

dt_default = DecisionTreeRegressor(random_state=42)

dt_default.fit(X_train_scaled, y_train)
```

```
# Evaluate the model on the training set

y_train_pred_dt_default = dt_default.predict(X_train_scaled)

rmse_train_dt_default      =      np.sqrt(mean_squared_error(y_train,
y_train_pred_dt_default))

mae_train_dt_default       =      mean_absolute_error(y_train,
y_train_pred_dt_default)
```

```
# Evaluate the model on the test set

y_test_pred_dt_default = dt_default.predict(X_test_scaled)

rmse_test_dt_default      =      np.sqrt(mean_squared_error(y_test,
y_test_pred_dt_default))

mae_test_dt_default = mean_absolute_error(y_test, y_test_pred_dt_default)


rmse_train_dt_default,      mae_train_dt_default,      rmse_test_dt_default,
mae_test_dt_default

(0.0, 0.0, 1.6722004164014066, 0.8692343374013606)
```

```
from sklearn.tree import export_graphviz

import graphviz
```

```
# Export the decision tree to a dot file

export_graphviz(dt_simplified,                        out_file="tree.dot",
feature_names=X_train.columns,                      filled=True,      rounded=True,
special_characters=True)
```

```
# Use Graphviz to read the dot file and create the visualization

with open("tree.dot") as f:

    dot_graph = f.read()

graphviz.Source(dot_graph)
```

```
# Train a simplified Decision Tree Regressor with limited depth

dt_simplified = DecisionTreeRegressor(random_state=42, max_depth=3)

dt_simplified.fit(X_train_scaled, y_train)
```

```
# Set the size of the plot

plt.figure(figsize=(20,10))
```



```
# Plot the simplified decision tree
plot_tree(dt_simplified,      filled=True,      feature_names=X_train.columns,
rounded=True, fontsize=10)
```

```
# Show the plot
plt.show()
```

Predictive Insights:

```
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeRegressor
from sklearn.preprocessing import StandardScaler
```

```
# Assuming previous feature selection identified relevant features, select
a subset for demonstration

selected_features = data.columns[3:] # Excluding 'Company', 'Time', and
'Financial Distress' for this example
```

```
# Preparing the data

X = data[selected_features]
y = data['Financial Distress']
```

```
# Splitting the data into training and testing sets

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42)
```

```
# Scaling the features

scaler = StandardScaler()

X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

```
# Training the Decision Tree Regressor

dt_regressor = DecisionTreeRegressor(random_state=42)

dt_regressor.fit(X_train_scaled, y_train)
```

```
# Extracting feature importances

feature_importances = dt_regressor.feature_importances_
```


```
# Mapping feature names to their importances
importance_dict = dict(zip(selected_features, feature_importances))
sorted_importance = sorted(importance_dict.items(), key=lambda item:
item[1], reverse=True)
```

```
# Displaying top 10 important features
```

```
sorted_importance[:10]
```

```
[('x48', 0.3530956354498967),
 ('x5', 0.3356245499264686),
 ('x81', 0.1798508200534849),
 ('x2', 0.01332075987668069),
 ('x73', 0.011058099374237991),
 ('x9', 0.0089512649925919),
 ('x19', 0.00582038706720564),
 ('x57', 0.00565021361675694),
 ('x36', 0.004707839913060336),
 ('x4', 0.004141671854292537)]
```

Forecasting and Scenario Analysis:



Model Deployment and Presentation:

```
from sklearn.tree import DecisionTreeRegressor
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
import numpy as np
import matplotlib.pyplot as plt
```

```
# Preparing the data

X = data.drop(columns=['Company', 'Time', 'Financial Distress'])
y = data['Financial Distress']
```

```
# Splitting the data into training and "future" data (simulating unseen
future data)

X_train, X_future, y_train, y_future = train_test_split(X, y, test_size=0.2,
random_state=42)
```

```
# Scaling the features

scaler = StandardScaler()

X_train_scaled = scaler.fit_transform(X_train)
X_future_scaled = scaler.transform(X_future)
```

```
# Training the Decision Tree Regressor

dt_regressor = DecisionTreeRegressor(random_state=42)
dt_regressor.fit(X_train_scaled, y_train)
```

```
# Predicting future business outcomes using the trained model

future_predictions = dt_regressor.predict(X_future_scaled)
```

```
# Scenario Analysis: Adjusting a key financial indicator (e.g., x5) by ±10%

X_future_scenario_increase = X_future_scaled.copy()

X_future_scenario_increase[:, X.columns.get_loc('x5')] *= 1.1 # Increase x5
by 10%

predictions_scenario_increase =
dt_regressor.predict(X_future_scenario_increase)

X_future_scenario_decrease = X_future_scaled.copy()
```

```
X_future_scenario_decrease[:, X.columns.get_loc('x5')] *= 0.9 # Decrease x5
by 10%

predictions_scenario_decrease =
dt_regressor.predict(X_future_scenario_decrease)
```

```
# Visualizing the Predictions for Different Scenarios

plt.figure(figsize=(14, 6))
```

```
# Original predictions

plt.subplot(1, 3, 1)

plt.hist(future_predictions, bins=30, color='gray', alpha=0.7)

plt.title('Original Predictions')

plt.xlabel('Financial Distress Score')

plt.ylabel('Frequency')
```

```
# Predictions with increased x5

plt.subplot(1, 3, 2)

plt.hist(predictions_scenario_increase, bins=30, color='blue', alpha=0.7)

plt.title('Predictions with Increased x5')

plt.xlabel('Financial Distress Score')
```

```
# Predictions with decreased x5

plt.subplot(1, 3, 3)

plt.hist(predictions_scenario_decrease, bins=30, color='red', alpha=0.7)

plt.title('Predictions with Decreased x5')

plt.xlabel('Financial Distress Score')

plt.tight_layout()

plt.show()
```